

Gradients Not Cleared before Backward Propagation

Context

In PyTorch, `optimizer.zero_grad()` **clears the old gradients** from last step, `loss_fn.backward()` **does the back propagation**, and `optimizer.step()` **performs weight update** using the gradients.

Problem

If `optimizer.zero_grad()` **is not used before** `loss_fn.backward()`, the **gradients will be accumulated from all `loss_fn.backward()` calls** and it will lead to the **gradient explosion**, which fails the training.

Solution

Developers **should not forget to use `optimizer.zero_grad()` before `loss_fn.backward()`.**

Existing Stage	Effect
Model Training	Error-prone

Example

```
Python
# PyTorch
# 3. Define a Loss function and optimizer  [...]
# 4. Train the network
for epoch in range(2): # training loop multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
+       # zero the parameter gradients
+       optimizer.zero_grad()
        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

